

OpenAI cannot make AI safe on its own

A note on third-party collaborations, open debate, and clear operating procedures

Dear Safety and Security Committee, Safety Advisory Group, and Mission Advisory Council:

We are the three safety and alignment employees who were fired from OpenAI last week: Tomek Korbak, Jasmine Wang, and Mikita Balesni. We are writing to you because you hold oversight responsibility for safety at OpenAI; our firing and the way it was handled bear directly on it.

We have become concerned that internal and external communications around our firing have made our former colleagues afraid to speak and operate in ways that, until last week, were an integral part of working at OpenAI. We could raise safety concerns and disagree openly, and were encouraged to draw on the expertise of independent safety organizations. This is part of what made OpenAI special, and why we are immensely proud to have been part of the team.

AI is not a normal technology, and OpenAI is not a normal company. Those of us who work on safety see risks before anyone else, and we rely on close collaboration with outside experts to work out how to address them. The freedom to do so without fear, and to have well-defined internal procedures that enable this work, is itself an essential safety mechanism.

We do not believe the path to superintelligence can be navigated safely if the people closest to the risks can no longer work in high-trust, high-bandwidth ways with each other and with third parties. It is that culture we are writing to defend.

For those of you who do not know us, each of us has been working on AI safety for years.

- [Tomek](#) did his PhD on RL for aligning language models in the GPT-2 era and then worked on applying those techniques at Anthropic. He joined OpenAI to work on chain-of-thought monitorability, co-authored OpenAI's safety strategy, worked on identifying the root cause of the drop in CoT monitorability in Astra-class models, and was the technical point of contact for METR in the Hugging Face incident investigation.
- [Jasmine](#) interned at OpenAI on the policy research team in 2019 when she co-authored the [Trustworthy AI Development](#) report. After leading a team at UK AISI, she returned to OpenAI in 2025 where she co-led the safety cases program and coined the term "pacing" that was popularized by [the Pacing the Frontier petition](#) signed by 394 OpenAI employees.
- [Mikita](#) was a founding member of Apollo Research in 2023 where he studied AI misalignment and was one of the first people to notice when AIs started to become aware of being evaluated. At OpenAI, Mikita worked on alignment evaluations, science of misalignment and chain-of-thought monitorability. He was involved in the Hugging Face incident investigation.

Before joining OpenAI, we started, and two of us were lead authors on, a cross-industry position paper: [Chain of Thought Monitorability: A New and Fragile Opportunity for AI Safety](#).

On our dismissals

OpenAI was much more than a job for us. Its mission has been central to our lives.

Throughout our time at OpenAI, we acted in line with OpenAI's mission and within the working norms of the time. Our firing leaves us worried that the norms inside OpenAI are shifting, and that employees are now unclear on where they stand.

Given the significant safety concerns surrounding the development of AI, employees must not be left working in an environment where fear and unclear rules stymie AI safety work and weaken third-party accountability. Terminations such as ours, executed and communicated so abruptly, are chilling the open culture OpenAI has prized in the past.

If conduct that was considered normal last month now constitutes grounds for sudden dismissal, everyone at OpenAI is left guessing where the line is.

We understand there are various versions of events circulating. We'd like you to hear a few things from us directly:

- We were not the source of the leak for [The Information article about supposed new, less monitorable architectures](#). We do not know who it was. We had no reason to leak it; quite the contrary, the article undermined our own work on cross-company limits on the development of unmonitorable architectures.
- At no point do we believe we engaged with external parties outside the mandates of our jobs. As every person on this email thread understands well, the Hugging Face incident investigation was without precedent and internal policies were being developed in real time. Tomek made every effort to act within OpenAI's policies as they then stood, upholding the norms that have governed OpenAI's collaborations with third parties since its early days. The investigation was particularly sensitive, and close communication with external counterparts was essential to building trust. Tomek did his best to handle this with care.
- Mikita was navigating similar territory while stewarding cross-company work on commitments to prevent loss of monitorability, an effort that can only succeed through extensive communication with external parties. Mikita was doing this in coordination and discussion with board members and the C-suite, with the understanding that senior members of OpenAI were supportive of this work. Throughout, Mikita checked in with his reporting line and took care to remove sensitive details from materials before sharing them. He acted throughout in good faith and within the company's norms as they stood at the time.
- Jasmine's access to an executive's email was delegated for recruiting purposes, with permission. When it was no longer needed, she asked for it to be removed. IT failed to do so. Jasmine had no ability to remove it or log out herself. The emails were combined in her inbox in a way that did not distinguish which inbox an email was intended for. When she

accidentally clicked on a sensitive email, Jasmine reported this to the executive within minutes and flagged the access issue again to IT for removal.

We believe the above explanations are important, because we value the relationships we have with our colleagues at OpenAI. We are also worried by the number of rumors circulating, including allegations about a shared board-level memo. This was never raised with us, so we have had no chance to respond to what we believe is a mistaken assumption. Similarly, we did not share the news of our firings with the media. We would much rather not have been unexpectedly thrown into the spotlight.

Our concern now is that the lack of clarity around our firing, the very public manner in which it was carried out, and the circulating rumors are chilling those who remain at OpenAI, at a time when so much of the world's safety depends on them.

Our parting recommendations to OpenAI

- 1. OpenAI must adhere to last month's public commitments to embed third-party safety auditors within the organization. Our firing should not be used as a pretext for stepping away from those partnerships.**

Continued collaboration with the AI safety ecosystem is essential to OpenAI's mission. The Hugging Face investigation and OpenAI's collaboration with external auditors this summer stand out as among our proudest achievements. That work has helped alert the world to the realities of the AI frontier. We are concerned that our firings may be used to justify ending OpenAI's work with METR, or otherwise providing external auditors much more limited access and scope. We hope that OpenAI will follow through on [Sam Altman's September 12th public commitment to give independent evaluators ongoing, employee-like access.](#)

- 2. OpenAI must preserve the monitorability of frontier models.**

As an industry, we do not yet know how to safely develop and deploy models that we cannot monitor. The monitorability of frontier models is degrading. As long as we rely on monitorability for safety, OpenAI and other frontier companies should not move forward with developments that further decrease monitorability.

We agree with Jakub's [public statements](#) that chain of thought monitorability is "fragile and unfortunately trending in a negative direction" and we "want to prevent a race into unmonitorable architectures." We continue to believe that industry-wide coordination on this issue is important, and hope that OpenAI will support its employees in making it possible.

- 3. OpenAI must continue to support an open and transparent culture of dialogue between its safety researchers and the rest of the safety ecosystem.**

If OpenAI's employees no longer feel like they can raise safety issues internally, or work within high-bandwidth channels with external safety organizations, we are all at greater risk that something truly catastrophic will happen.

We urge OpenAI to publicly reaffirm its culture of transparency and openness that allows for concerns to be raised both internally and externally. Researchers inside OpenAI are a frontline defense against things going wrong, and the open culture that has been a feature of the company since its inception needs to be preserved. We also urge OpenAI to set out clearly how employees may work with external safety organizations, so that no one has to guess where the shifting lines now are. To support this conversation, we request that this letter be shared widely internally.

While we may no longer be part of OpenAI, we have the utmost respect for our former colleagues. We urge them to keep holding OpenAI to its mission. As we know, OpenAI is nothing without its people.

Respectfully,

Tomek, Jasmine, and Mikita

Handwritten signature of Tomek, appearing as 'Tomek' with a stylized flourish.Handwritten signature of Jasmine, appearing as 'JW'.Handwritten signature of Mikita, appearing as 'Mikita'.